



Nuzzle × Gigalogy MAIRA AI

Presentation Script

Complete Slide-by-Slide Presenter Speech & Executive Companion • 15 Widescreen Slides

15-Slide Master
Deck

Estimated Speech Time: 15–18
Minutes

Target Audience: Judges, Investors & Technical
Evaluators

How to use this script: This guide provides a word-for-word spoken script written in clear, energetic, and natural English. Each slide card contains what is shown on screen, the exact spoken words to say, key bullet points to reference if you prefer summarizing, and a smooth transition sentence into the next slide.

Gigalogy MAIRA AI Integration in Nuzzle Pet Social Media

Goal: Welcome the audience, establish your identity, and set the high-stakes narrative.

Audience Sees: Dark executive title slide, Nuzzle logo, 5 live metric cards (4 Live Profiles, 3 Vector Indices, GPT-4o Vision, 2ms Poison Guardrail, 100% Uptime), and deployment status.

VERBATIM PRESENTER SPEECH

"Good day, everyone! Welcome to our technical walkthrough of the **Nuzzle AI Platform** and our deep production integration with **Gigalogy MAIRA Managed RAG**."

"At its heart, Nuzzle is Bangladesh's premier pet social and animal welfare platform. But pets cannot speak. When a dog accidentally ingests chocolate at 2:00 AM, or when an indoor cat slips past the front door in a bustling city like Dhaka, seconds literally mean the difference between life and death."

"Generic AI chatbots often fail here because they hallucinate medical dosages or lack local context. Over the next fifteen minutes, we're going to walk you through how we solved this by deploying **four specialized Gigalogy MAIRA AI profiles**, trained on proprietary high-dimensional pet datasets, backed by a **sub-two-millisecond deterministic emergency guardrail**, and wrapped in an ultra-resilient zero-downtime architecture."

KEY TALKING POINTS TO REMEMBER:

- Nuzzle integrates social community + clinical veterinary triage + visual pet recovery.
- 4 live isolated MAIRA profiles with custom system prompts and vector datasets.
- Production-ready on Next.js 16 Edge, Vue 3, and Supabase pgvector.

Bridge to Next Slide: "To appreciate why this architecture matters, let's first examine why traditional AI approaches fail pet parents in real emergencies."

Strategic Motivation: Why Gigalogy MAIRA?

Goal: Explain the core limitations of generic LLMs and show why isolated domain profiles win.

Audience Sees: Comparison showing Generic Large LLMs (Prompt pollution, hallucinations, high token costs) vs. Gigalogy MAIRA Advantage (Dedicated system profiles, sub-500ms retrieval, zero hallucinations).

VERBATIM PRESENTER SPEECH

"When engineers first add AI to an app, the common temptation is to just send every prompt to ChatGPT. But in veterinary medicine, that approach has three critical flaws."

"First is **prompt pollution**. When you force one single prompt to act as an emergency vet, a lost-and-found investigator, a local clinic directory, and a pet diet advisor, it gets confused and gives bland, generic responses."

"Second is **hallucination**. An ungrounded LLM might fabricate a toxic dosage or recommend a clinic in New York instead of Dhanmondi, Dhaka."

"And third is **token cost and latency**. Stuffing large clinic databases into prompt context wastes tokens and slows down emergency responses."

"Gigalogy MAIRA solves this elegantly. Instead of one bloated prompt, MAIRA gives us **isolated, specialized domain profiles**. Each profile connects directly to its own trained vector knowledge base. This guarantees sub-500-millisecond retrieval, razor-sharp domain accuracy, and zero medical hallucinations."

KEY TALKING POINTS TO REMEMBER:

- Generic LLMs suffer from context dilution, prompt pollution, and geographic hallucinations.
- MAIRA decouples domain responsibilities into isolated, UUID-bound intelligence nodes.
- Every answer is strictly grounded in verified clinical guidelines and local Dhaka data.

Bridge to Next Slide: "Now, let's look at how our full-stack engineering connects our frontend users to this intelligence cluster."

Full-Stack Architecture & API Ingress Pipeline

Goal: Walk through the request journey and introduce our triple-tier failover strategy.

Audience Sees: 4 horizontal pipeline blocks (Client Ingress → Edge Gateway → MAIRA Intelligence Cluster → Persistence Tier) plus the bottom Triple-Tier Failover Engine (MAIRA → GPT-4o → Offline Heuristics).

VERBATIM PRESENTER SPEECH

"Here you see our end-to-end request pipeline. It flows seamlessly across four decoupled stages."

"It starts on the left with our **Vue 3 Single Page App**. When a user snaps a photo of a lost pet or chats with our AI vet, the request hits our **Next.js 16 Edge API Gateway**."

"Here at the gateway, every request is validated with strict Zod schemas, authenticated via Supabase JWT, and rate-limited to 60 requests per minute to prevent scrapers."

"Next, our router determines where the query belongs. If it's a symptom inquiry, it queries MAIRA's PawDoctor profile. If it's a sighting photo, it triggers our multimodal vision pipeline."

"Finally, at the bottom, notice our **Triple-Tier Failover Strategy**. Tier 1 is always Gigalogy MAIRA. But if MAIRA ever experiences a cloud disruption or rate limit, our router automatically fails over to Tier 2—an OpenAI GPT-4o-mini cluster. And if the entire internet drops, Tier 3 kicks in: local in-memory deterministic heuristic tables. Our users never, ever see a 500 server crash."

KEY TALKING POINTS TO REMEMBER:

- Client Layer (Vue 3) → Edge Gateway (Next.js 16) → AI Core (MAIRA) → DB (Supabase).
- Zod schema safety, JWT session verification, token-bucket rate limiting.
- Zero-crash guarantee via 3-tier active failover hierarchy.

Bridge to Next Slide: "To see how all these services map out across our infrastructure, let's look at our complete system blueprint."

System Blueprint & Four-Tier Architectural Topology

Goal: Provide a structural blueprint showing client, gateway, multi-model AI, and vector database.

Audience Sees: 4 distinct enterprise architectural columns: Tier 01 Client & UI Layer, Tier 02 Edge API Gateway, Tier 03 Clinical AI & Vision Core, and Tier 04 Data & Vectors Tier with Supabase cloud.

VERBATIM PRESENTER SPEECH

"This blueprint illustrates our clean four-tier architectural separation."

"**Tier 1 is our Client Layer.** It provides intuitive interfaces: camera capture for rescue sightings, reactive chat for veterinary triage, and a live geographic clinic radar."

"**Tier 2 is our Edge Gateway.** It acts as the traffic controller. It normalizes inputs, verifies user permissions, and ensures secure, authenticated communication."

"**Tier 3 is our Multimodal Intelligence Cluster.** Here, we don't rely on just one model. We combine a sub-two-millisecond deterministic emergency verifier, Hugging Face's Qwen2.5-VL vision model for biometric pet facial recognition, and Gigalogy MAIRA RAG for deep veterinary knowledge."

"And **Tier 4 is our Persistence Layer** hosted on Supabase Cloud. It leverages PostgreSQL with the pgvector extension for sub-15ms cosine distance matching, S3-compatible storage for high-res sighting photos, and WebSockets for real-time rescue alerts."

KEY TALKING POINTS TO REMEMBER:

- Decoupled microservices architecture: Frontend, Gateway, AI Models, Storage.
- Hybrid AI: combines rule-based deterministic guards with neural vision and vector RAG.
- Hosted on robust cloud infrastructure in the Asia-Pacific region.

Bridge to Next Slide: "Now, let's zoom in specifically on the AI pipeline to see how data flows between our models in real time."

Enterprise AI Integration Architecture Diagram

Goal: Showcase the high-resolution AI data flow diagram and explain the 5-stage inference cascade.

Audience Sees: High-res visual diagram showing Input Ingestion → Stage 1 Guardrails (<2ms) → Stage 2 Vision Biometrics (~800ms) → Stage 3 MAIRA RAG (~450ms) → Stage 4 Fallback → Stage 5 pgvector Broadcast, flanked by 3 key KPI cards.

VERBATIM PRESENTER SPEECH

"On this slide, you see our complete **Enterprise AI Integration Pipeline**."

"When a pet parent or animal rescuer submits data—whether it's a photo, symptom description, or GPS coordinate—it passes through a synchronized five-stage cascade."

"In **Stage 1**, our sub-two-millisecond deterministic guardrail acts as our first line of defense. If someone accidentally uploads a photo of a coffee cup or a human selfie, it's rejected immediately without wasting expensive GPU cycles. Even more crucially, if someone types 'my dog just ate chocolate,' it intercepts the query in less than 2 milliseconds with life-saving emergency actions."

"In **Stage 2**, sighting photos are dispatched to our Hugging Face vision router running Qwen2.5-VL-72B to extract anatomical biometrics—like ear geometry and coat markings."

"In **Stage 3**, textual inquiries hit Gigalogy MAIRA's vector RAG engine in approximately 450 milliseconds."

"And in **Stage 5**, verified matches are indexed into Supabase pgvector and broadcast instantly across our rescuer network via WebSockets."

KEY TALKING POINTS TO REMEMBER:

- Stage 1 Guardrail: <2ms latency, filters invalid media and catches lethal toxins immediately.
- Stage 2 Vision AI: Qwen2.5-VL-72B biometric feature extractor on Hugging Face router.
- Stage 3 MAIRA: 4 UUID profiles delivering grounded domain answers in ~450ms.

Bridge to Next Slide: "Now let's explore our first major clinical pillar: PawDoctor AI."

Pillar 1: PawDoctor AI — Clinical Triage & Poison Fast-Path

Goal: Demonstrate life-saving veterinary triage and deterministic poison guardrails.

Audience Sees: Left: Deterministic Emergency Guardrail card with lethal thresholds (Chocolate >20mg/kg, Lilies feline toxicity, Paracetamol, Grapes). Right: MAIRA Clinical Triage Profile card and sample clinical response.

VERBATIM PRESENTER SPEECH

"Slide 6 showcases **PawDoctor AI**, our 24/7 veterinary triage companion."

"PawDoctor is powered by MAIRA Profile ID 0a8fd1e8. But we designed it with a crucial hybrid safety mechanism."

"Look at the left column: our **Deterministic Poison Fast-Path**. Medical emergencies cannot wait for a cloud LLM to think. Substances like dark chocolate contain theobromine, which is lethal to canines above 20 milligrams per kilogram. Lilies cause acute, fatal renal failure in cats with just a single speck of pollen. Paracetamol destroys feline red blood cells."

"Our edge guardrail scans symptoms against strict clinical tables. In under two milliseconds, it detects dangerous toxins and delivers bold, step-by-step first-aid instructions alongside the phone number of the nearest 24/7 emergency hospital."

"For non-acute queries, MAIRA steps in with warm, empathetic clinical guidance—asking clarifying questions about hydration, gum color, and lethargy, always reinforcing that it is an assistive triage tool and not a replacement for an in-person veterinarian."

KEY TALKING POINTS TO REMEMBER:

- Sub-millisecond poison calculation intercepts chocolate, lily, grape, and drug poisonings.
- Zero hallucination of dangerous medical dosages.
- Combines deterministic emergency safety with empathetic, conversational triage.

Bridge to Next Slide: "Our second pillar addresses every pet parent's worst nightmare: losing a beloved pet."

Pillar 2: Lost & Found Pet Safety Radar

Goal: Explain multimodal visual biometric pet matching and geospatial correlation.

Audience Sees: Left: Multimodal Search Architecture card (Camera intake, candidate retrieval, Qwen2.5-VL biometric scoring). Right: Real benchmark case study (Milo the Golden Retriever in Gulshan-2, matched at 0.92 confidence).

VERBATIM PRESENTER SPEECH

"Losing a pet is terrifying. In a dense city like Dhaka with narrow alleyways and heavy traffic, finding a lost dog or cat is like searching for a needle in a haystack."

"Traditional lost pet boards rely purely on text descriptions like 'brown dog with white paws.' That's far too vague."

"With our **Lost & Found Radar**, powered by MAIRA Profile ID a10e6172, anyone who spots a wandering animal can simply take a photo with their phone. Our pipeline extracts anatomical biometric markers: facial masks, coat patterns, ear posture, and collar tags."

"Look at the benchmark case study on the right. When 'Milo,' a lost Golden Retriever in Gulshan-2, was sighted by a volunteer, our system compared the photo against 50 registered lost pets. It identified a 0.92 biometric correlation in under two seconds, correctly confirming Milo's white chest patch and red reflective collar, and immediately alerted the owner."

KEY TALKING POINTS TO REMEMBER:

- Combines visual biometric comparison with geospatial Haversine distance filtering.
- Evaluates coat patterns, facial markings, and collar tags with 0–100% confidence.
- Validated against a 50-pet benchmark dataset across Dhaka districts.

Bridge to Next Slide: "Let's take a look at what this experience actually looks like in our live production application."

Pillar 2 • Live Application Demo: Visual Biometric Scanner

Goal: Show live UI screenshots and prove the user workflow is intuitive and production-ready.

Audience Sees: Left: Real screenshot of the Nuzzle AI Visual Radar Scanner showing sighting photo preview, biometric match score (92%), and contact owner button. Right: 3 UI feature breakdown cards.

VERBATIM PRESENTER SPEECH

"Here on Slide 8 is an actual screenshot of our live production interface in Vue 3."

"Notice how clean and intuitive the user experience is. A user simply taps 'Upload Sighting Photo' or uses their phone camera. The image is compressed and previewed in real time."

"Our backend endpoint `/api/lost-found/ai-match` runs the sighting through our Hugging Face Qwen2.5-VL-72B model in approximately 1.8 seconds."

"The UI presents the matched pet side-by-side with the sighting photo, displays the exact 92% match score, and lists the AI's reasoning. And most importantly, it gives the finder a one-tap phone dialer to call the owner or dispatch rescue coordinates. No complex signups, no barriers—just fast, life-saving reunification."

KEY TALKING POINTS TO REMEMBER:

- Mobile-optimized responsive Vue 3 SPA interface.
- Side-by-side visual photo comparison with clear confidence rating.
- One-tap phone calling and location dispatch for instant rescue action.

Bridge to Next Slide: "When a pet is found injured or sick, what happens next? That brings us to Pillar 3: our Specialist Clinic Navigator."

Pillar 3: Vet Clinic & Telemedicine Specialist Navigator

Goal: Explain geographic clinic routing, avian/exotic filters, and direct appointment booking.

Audience Sees: Left: Clinical Routing Architecture card with verified Dhaka clinics (Cascade 24/7 in Dhanmondi, Gulshan Pet Care, Uttara Companion, Mirpur Mobile). Right: Telemedicine Booking & Specialist Routing card with appointment slots.

VERBATIM PRESENTER SPEECH

"Finding a veterinary clinic in an emergency is often frustrating. Many general clinics do not have orthopedic surgery equipment, avian veterinarians, or 24/7 night staff."

"Our third pillar, powered by MAIRA Profile ID 9884416a, acts as an intelligent clinical routing navigator."

"We indexed verified veterinary facilities across all key districts of Dhaka—including Dhanmondi, Gulshan, Uttara, and Mirpur. Each entry has metadata tags for 24/7 emergency care, ICU incubators, radiography, ultrasound, and specialized species coverage like birds and exotic pets."

"When a pet parent reports 'my parrot has a broken wing' or 'my dog needs emergency surgery,' MAIRA evaluates the exact geographic proximity and clinical specialty. It filters out clinics without proper equipment and returns only the verified facilities capable of handling that specific emergency, complete with direct appointment booking and phone numbers."

KEY TALKING POINTS TO REMEMBER:

- Verified database across Dhaka: Dhanmondi, Gulshan, Uttara, Mirpur.
- Clinical capability tags: 24/7 ICU, digital X-ray, orthopedic surgery, avian/exotic care.
- Instant telemedicine booking and telephone hotline integration.

Bridge to Next Slide: "Long-term pet wellness also depends on daily care. Let's look at Pillar 4: AI Nutrition and Marketplace."

Pillar 4: Marketplace AI Nutrition & Dietary Formulation

Goal: Show personalized dietary formulas, allergen exclusion, and genuine marketplace conversions.

Audience Sees: Left: Dietary Advisory Engine card (Caloric formulation, life-stage profiling, zero-poultry/wheat-free filters). Right: Product Retrieval & Conversion Funnel showing genuine items (Royal Canin Gastrointestinal, Taste of the Wild, Virbac Gel).

VERBATIM PRESENTER SPEECH

"Our fourth pillar connects healthcare intelligence with genuine pet commerce."

"Counterfeit pet food and inappropriate diets are a huge problem in developing markets. Feeding high-protein adult food to a puppy with a sensitive stomach can cause severe gastrointestinal distress."

"Our **Marketplace AI Nutritionist** is powered by MAIRA Profile ID 66e15325. It calculates precise Resting Energy Requirements (RER) and Maintenance Energy Requirements (MER) based on the pet's breed, age, neutered status, and weight."

"It cross-references dietary requirements against our verified catalog of genuine products—brands like Royal Canin, Taste of the Wild, and Virbac. If a pet parent says 'my Persian cat is allergic to chicken and vomits dry kibble,' MAIRA excludes poultry byproducts and recommends clinically verified grain-free or novel protein diets with accurate pricing in Bangladeshi Taka."

KEY TALKING POINTS TO REMEMBER:

- Caloric calculation based on clinical RER/MER formulas.
- Allergen exclusion (wheat-free, zero-poultry, renal support).
- Guaranteed genuine inventory with local pricing in Bangladeshi Taka (BDT).

Bridge to Next Slide: "Now that we've seen all four pillars in action, let's look at how we engineered and indexed the underlying data."

Dataset Engineering: Curating & Training MAIRA Vector Knowledge

Goal: Pull back the curtain on the data engineering, schema design, and vector ingestion.

Audience Sees: 3 Dataset Detail Cards: Lost & Found Pet Radar Dataset (50 benchmark records), Veterinary Clinics Directory (16 certified Dhaka hospitals), and Pet Nutrition & Marketplace Catalog (25 genuine formulas).

VERBATIM PRESENTER SPEECH

"An AI model is only as intelligent as the data it retrieves. On Slide 11, we pull back the curtain on our **Dataset Engineering Pipeline**."

"We curated, cleaned, and indexed three high-dimensional datasets into Gigalogy MAIRA's cloud."

"First, our **Lost & Found Benchmark Dataset** containing 50 diverse pet records across Dhaka, rich with GPS coordinates, distinctive marks, and image URLs."

"Second, our **Veterinary Clinic Directory** containing 16 certified animal hospitals across 4 Dhaka zones, complete with ICU availability, emergency contact numbers, and doctor specialties."

"And third, our **Marketplace Catalog** with 25 genuine prescription formulas, caloric densities, and ingredient transparency."

"During ingestion, we engineered custom key-value serialization to ensure MAIRA's vector indexers capture both semantic meaning and exact factual attributes without loss."

KEY TALKING POINTS TO REMEMBER:

- 3 curated datasets: 50 benchmark pets, 16 verified clinics, 25 nutrition formulas.
- Structured key-value serialization for high vector indexing precision.
- Clean separation of datasets aligned with the 4 isolated MAIRA profiles.

Bridge to Next Slide: "Next, let's explore the technical hurdles our team solved during API implementation."

Technical Implementation: API Protocols & Header Security

Goal: Demonstrate real engineering competence, resolving Gigalogy auth quirks, and payload optimization.

Audience Sees: 4 Implementation Cards: Authentication Discovery (Bearer vs. Raw API Key), Strict Prompt Synthesis, Payload Optimization (<10MB Base64), and Edge Deployment Pipeline.

VERBATIM PRESENTER SPEECH

"Real-world integration always uncovers technical nuances. On Slide 12, we highlight key engineering discoveries our team resolved."

"First was **Header Security and Authentication**. While standard REST APIs use `Authorization: Bearer`, Gigalogy MAIRA expects the raw API key directly in the header without the 'Bearer' prefix. We encapsulated this discovery into a hardened, reusable singleton client `gigalogyClient.ts`, preventing silent 401 authentication drops."

"Second was **Payload Optimization**. Mobile cameras capture photos exceeding 15 megabytes. We implemented client-side canvas compression, guaranteeing all photo payloads remain under 10 megabytes with sub-300ms serialization times."

"And third, all environment keys are securely injected at build-time with strict TypeScript interfaces, ensuring zero leaked credentials."

KEY TALKING POINTS TO REMEMBER:

- Resolved raw header vs. Bearer prefix requirement in Gigalogy MAIRA API.
- Client-side image downsampling prevents mobile network timeouts.
- TypeScript type safety across all request and response schemas.

Bridge to Next Slide: "How do we guarantee that our system remains bulletproof even during cloud outages? That's Slide 13."

Resilience & Security: Zero-Downtime Resilience Engine

Goal: Detail circuit breaking, rate limiting, and local Supabase session security.

Audience Sees: 3 Cards: Triple-Tier Failover Engine (Active circuit breaker), Token-Bucket Rate Limiter (60 req/min per IP), and Supabase Auth & JWT Session Security.

VERBATIM PRESENTER SPEECH

"In production, cloud services can and do fail. On Slide 13, we showcase our **Zero-Downtime Resilience Engine**."

"Our gateway implements an automated **Circuit Breaker**. If the MAIRA API returns three consecutive timeouts or a 429 rate limit, the circuit trips immediately. Traffic gracefully switches to Tier 2 (our OpenAI GPT-4o-mini cluster) in under 50 milliseconds, with zero disruption to the user."

"To prevent abuse and bot spam, our in-memory token-bucket rate limiter restricts traffic to 60 requests per minute per IP address."

"And all pet claims, medical records, and user chats are guarded by Supabase Auth with Row Level Security (RLS) policies, ensuring pet owners only access their own private medical consultations."

KEY TALKING POINTS TO REMEMBER:

- Automated circuit breaker trips on 3 consecutive errors or rate limits.
- Sub-50ms seamless failover between MAIRA, OpenAI, and local heuristics.
- Token-bucket rate limiting and Supabase Row Level Security (RLS).

Bridge to Next Slide: "Let's look at the hard numbers and automated telemetry verifying this performance."

Production Benchmarks & Live Verification Telemetry

Goal: Provide concrete quantitative proof of speed, uptime, and test suite verification.

Audience Sees: Left: 4 Live Benchmark Metric Boxes (450ms MAIRA RAG, <2ms Poison Guard, 92% Vision Biometrics, 100% Uptime). Right: Verification Telemetry Log showing all test suites passing with 200 OK.

VERBATIM PRESENTER SPEECH

"Engineers trust data, not promises. On Slide 14, we present our **Production Telemetry and Live Verification Benchmarks**."

"Every single component has been rigorously stress-tested."

"Our average MAIRA RAG retrieval latency clocks in at **450 milliseconds**. Our emergency poison guardrail executes in **under 2 milliseconds**. Our multimodal vision biometrics achieve **92% top-match accuracy** across our benchmark dataset. And across all test cycles, we maintained **100% service uptime**."

"On the right is our live telemetry log. You can see all four core endpoints—PawDoctor triage, lost pet matching, clinic directory search, and health checks—returning HTTP 200 OK with clean payload verification. The platform is robust, responsive, and ready for scale."

KEY TALKING POINTS TO REMEMBER:

- 450ms average RAG response latency; <2ms poison safety guard.
- 92% biometric correlation accuracy on real pet sightings.
- 100% uptime with automated test suite validation.

Bridge to Next Slide: "Finally, let's wrap up with our future roadmap and vision for Nuzzle."

Future Roadmap, Production Conclusion & Q&A

Goal: Share exciting upcoming features, deliver a memorable closing, and invite questions.

Audience Sees: 3 Roadmap Phase Cards (Q3 2026 Edge Onnx Biometrics, Q4 2026 Automated Clinic Dispatch, Q1 2027 Bilingual Voice Triage) and a prominent Conclusion & Q&A Box.

VERBATIM PRESENTER SPEECH

"To conclude, our journey doesn't stop here. Looking forward across our roadmap:"

"In **Q3 2026**, we are deploying client-side ONNX vision models directly inside the browser for instant offline collar and breed detection."

"In **Q4 2026**, we are integrating automated WhatsApp and SMS dispatch to notify veterinary ambulances the instant an acute trauma is detected."

"And in **Q1 2027**, we are launching bilingual voice emergency triage, allowing panicking pet owners to verbally speak symptoms in Bengali or English during high-stress situations."

"In conclusion: by combining Gigalogy MAIRA's managed RAG profiles with deterministic safety guardrails and modern full-stack engineering, Nuzzle has transformed from a social network into a life-saving veterinary companion for thousands of animals across Bangladesh."

"Thank you so much for your time and attention! We would now love to open the floor to any questions."

KEY TALKING POINTS TO REMEMBER:

- Roadmap: On-device ONNX vision, WhatsApp emergency dispatch, voice triage in Bengali.
- Core achievement: Bridged social community with clinical veterinary AI.
- Confidently invite questions from the evaluators.

Evaluator Q&A Cheat-Sheet: Top 6 Anticipated Questions & Winning Answers

Be fully prepared for technical judges and business evaluators. Here are the exact answers to the most common questions:

Q1: "Why not just use OpenAI GPT-4 directly instead of Gigalogo MAIRA?"

"Great question. Generic GPT-4 is an ungrounded foundation model. When asked about pet poisonings or local clinics, it suffers from context drift, high latency, and severe hallucinations about local Dhaka geography. Gigalogo MAIRA provides managed, isolated RAG profiles. Each profile is pre-indexed with our verified clinical datasets, delivering sub-500ms answers that are 100% grounded in veterinary facts and local Bangladesh clinic data at a fraction of the token cost."

Q2: "What happens if a pet owner submits an acute emergency while Gigalogo MAIRA is offline?"

"Our system never fails during acute emergencies because of our deterministic guardrail. Severe toxins like dark chocolate, lilies, and paracetamol are caught in under 2 milliseconds by our local edge engine before any cloud API is even called. Furthermore, our triple-tier architecture automatically fails over to OpenAI GPT-4o-mini within 50 milliseconds if MAIRA is unreachable, backed by local heuristic directories as an offline failsafe."

Q3: "How accurate is your pet facial and coat biometric matching?"

"In our 50-pet benchmark evaluation across diverse canine and feline breeds in Dhaka, our multimodal vision comparator running Qwen2.5-VL-72B achieved a 92% top-match accuracy. It doesn't just look at color—it analyzes facial masks, muzzle shading, ear geometry, collar presence, and distinctive markings to differentiate between similar-looking pets."

Q4: "How do you prevent non-pet spam or inappropriate images from wasting compute?"

"Our Next.js edge gateway performs strict pre-inference validation. We enforce a 10MB payload limit, run base64 header verification, and execute lightweight heuristic classification. If a user uploads a human selfie, scenery, or an object, our API intercepts it with a 400 validation error in under 2 milliseconds, entirely preserving our AI rate limits."

Q5: "How are you handling bilingual Bengali and English language support?"

"We support Bengali natively through a hybrid approach. For emergency screening, our regex and heuristic filters recognize common Bengali toxic substance names (like 'চকলেট', 'লিলি ফুল', 'প্যারাসিটামল', 'বমি') with sub-5ms latency. For conversational triage, MAIRA processes Bengali syntax directly, recognizing colloquial Bangladeshi terms and replying in fluent, comforting Bengali with clear bulleted first-aid steps."

Q6: "What is your operational cost and mobile latency impact?"

"Our architecture is highly cost- and performance-optimized. By filtering acute toxic emergencies with our local 2ms guardrail, we completely avoid LLM token consumption on routine poison checks. For RAG queries, MAIRA caches vector embeddings and returns responses in 450ms. With client-side image compression, mobile payloads remain under 10MB even on 4G cellular connections across Bangladesh."